

Rohith Siliveri

Applied AI Engineer

rohith.career.ai@gmail.com | linkedin.com/in/rohith1411 | github.com/rsiliveri14 | rohithsiliveri.com

SUMMARY

Applied AI engineer focused on coding-agent evaluation, verifiers that cannot be gamed, and production machine learning in banking and insurance. Builds SWE-bench, Terminal-Bench, and Harbor-style tasks; runs blinded agent evals; ships claims, fraud, and document models.

EXPERIENCE

Applied AI Engineer - PNC Financial Services

Jul 2025 - Present | USA

- Own applied ML for banking -- risk, fraud, documents, and decision support -- with held-out evaluation and audit-ready runs.
- Build Harbor and Terminal-Bench-style agent evals: pinned Docker, oracles, and verifiers that fail unless the work is real.
- Trace agent failures to the model, the prompt, or the environment so production AI can be measured, not guessed.

AI Fellow / Reviewer / Specialist - Handshake AI

Jan 2026 - Aug 2026 | Remote | Concurrent with PNC

- Authored SWE-bench, Terminal-Bench, and Harbor tasks from real merged PRs -- pinned Docker, golden solutions, fail-to-pass tests.
- Ran blinded evals of coding agents (Claude Code, OpenCode, Grok CLI) and wrote line-cited pass/fail rationales.
- Produced visual bug-fixing and preference-ranking data so labs can train and score agents, not only chat models.

Expert AI Contributor - Snorkel AI

Mar 2026 - Aug 2026 | Remote | Concurrent with PNC

- Validated and repaired Harbor packages: hermetic Docker, no leakage, and instruction-test-solution alignment.
- Rewrote over-prescriptive prompts into real engineering asks and replaced shallow, gameable verifiers.
- Built OSWorld-style computer-use GUI tasks and GPU ML benchmarks in the MLE-bench / RE-bench family.

Graduate Student Researcher - Binghamton University

Jan 2025 - May 2025 | NY

- Faculty research: an end-to-end residential property valuation system, from housing data through a deployed app.
- Benchmarked five regressors; selected XGBoost after it beat a Linear Regression baseline of $R^2 \sim 0.69$.
- Served the model with FastAPI, PostgreSQL, Docker, AWS EC2/S3, and GitHub Actions CI/CD.

AI Engineer - Assurant

Jan 2022 - Jul 2023 | India

- Built claims-triage models for property and auto from claims, policy, and customer history; evaluated on precision, recall, and ROC-AUC.
- Shipped Airflow training pipelines with MLflow tracking and FastAPI services on Docker and Kubernetes.
- Extracted fields from messy claim forms with OCR, spaCy, and Hugging Face; trained XGBoost fraud models as a rare-class problem.

EDUCATION

M.S. Computer Science (AI) - SUNY Binghamton, NY

B.Tech Computer Science - Anurag University, India

SKILLS

Eval: SWE-bench, Terminal-Bench, Harbor, OSWorld, MLE-bench, pytest, oracle/nop scoring

ML: Generative AI, LLMs, Agentic AI, PyTorch, TensorFlow, Hugging Face, XGBoost, scikit-learn

Languages: Python, TypeScript, JavaScript, Go, Java, Rust, SQL, Shell

Infra: Docker, Kubernetes, FastAPI, GitHub Actions, Jenkins, AWS, PostgreSQL, Kafka, Airflow, MLflow